

Prepared release draft - activation and final checks pending

Real versus AI-generated image classification with an AI-positive score, local CPU inference, batch JSONL and a web interface. Results below are self-evaluated public-data measurements. Organizer hidden scores and qualifying rank are unknown.

Model, data and operating point

Frozen generic **CLIP ViT-L/14** image tower plus our trained logistic head. Training: **8,000 CIFAKE + 6,239 GenImage BigGAN/SD1.5** images; original and identical crop/JPEG views, 90% GenImage loss weight. C=10 selected on internal validation. The pretrained tower is credited to OpenAI; we trained the head.

Inference: EXIF orientation, official 224px CLIP centre crop, L2-normalized embedding, linear logit and temperature scaling. Temperature **0.946907**; AI threshold **0.821428**. Temperature uses separate calibration data; threshold uses internal validation, not external labels.

Calibration: 9,969 CIFAKE + 781 GenImage; validation sizes below. Exact/perceptual overlap screening and grouped splits are heuristic. Backbone pretraining overlap is unknown. GenImage-derived head: noncommercial research use.

Runnable interface and attribution

Shared CPU CLI/API; single image or batch with one loaded model. JPEG/PNG/WebP uploads up to **25 MiB / 40 MP**. EXIF is separate and does not alter predictions; C2PA is not checked.

Packaged image tower + head: **608.36 MB**, verified by SHA-256. CLIP inference requires torch and the application dependencies, with no CLIP/torchvision/ftfy/regex install. Measured CPU scoring: approximately **0.38 seconds/image**; fresh-download setup timing is pending.

CLIP explanation integration and its 40-image audit are pending. The planned input-gradient saliency and masking check describe model influence, not verified visual defects. Human usefulness review remains pending.

Measured development performance at the stated operating point

Evaluation	AUC	Macro-F1	Accuracy	Real FPR	AI recall
GenImage validation (783)	0.991	0.950	95.0%	1.9%	92.3%
CIFAKE validation subset (4,000)	0.910	0.759	76.7%	5.0%	58.4%
External dev - original	0.771	0.604	63.5%	11.7%	38.8%
External dev - format matched	0.789	0.608	64.2%	10.0%	38.5%
GLIDE/DALLE - second use	Pending; no score claimed				
GLIDE/DALLE matched - second use	Pending; no score claimed				

External rows average generator/source pairs (2,000 unique dev images per format). Development numbers use the training encoder. CPU dtype/export checks found zero label disagreements on measured samples, score drift up to 0.0097; parity evidence, not another accuracy claim. Organizer data and baseline were not supplied.

Published acceptance failure: balanced L/14 passed 11 of 12 internal checks, but as-distributed LDM/LAION real FPR was **22.0%**, versus v0.2.0's **12.8%** (allowed increase: 5 points). The gate **failed**. Release preparation prioritizes mean AUC improvement **0.647 to 0.771** original and **0.653 to 0.789** matched, accepting this disclosed tradeoff. Threshold retry and COCO augmentation did not fix the failed check.

Confusion matrix and evaluation integrity

GenImage validation; rows actual real/AI, columns predicted real/AI: **[[362, 7], [32, 382]]**.

354 COCO reserved real photographs remain unscored; a first-use false-positive check follows freeze. It measures real-photo FPR only.

GLIDE/DALLE was already evaluated for v0.2.0. Any new result is a disclosed second use after freeze, never described as a fresh blind test. Old release scores and weights remain unchanged.

Limitations and submission checks

False alarms and missed unseen AI images remain material. Scores are not guarantees of origin or event truth. Per-domain calibration may not transfer; centre cropping omits image borders. Private user images are unavailable and were not retested.

The robustness interface includes JPEG, resizing, blur and a clearly labelled screenshot simulation. **New-model robustness measurements are pending**; the old ResNet's results are not attributed to CLIP. Generator attribution and caption consistency are not implemented.